Epicure: A Meal Recognition System

Presentation by:

Farhan Hormasji

Max Ellison

Tarang Chugh

Agenda

- Overview
  - Objective
  - Datasets
- Proposed Algorithm
  - Preprocessing
  - Feature Extraction
  - Classifier Training
- Results and Critique
  - Accuracy
  - Discussion

Objective

- To develop a comprehensive and robust meal recognition system, able to identify at least 10 different classes of food
- To understand the interdependence and working of various vision modules: preprocessing, segmentation, feature extraction, etc










Dataset – I (Into the wild)

- **Training Set:** 50 images for each class
- Collected images from internet and few captured by us
- Manually filtered to ensure a mix bag of different view-angles and food-forms, illumination variations
- **Testing Set:** 100 images, at least 10 images for each class (25 images with multiple labels)




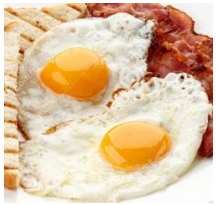


Dataset- II (Conducive Dataset)

- **Training Set:** 715 Images in 7 Classes
- Collected images from internet with easy to segment background
- To understand the impact of background noise on classification accuracy
- **Testing Set:** 70 images for 7 classes




Proposed Algorithm

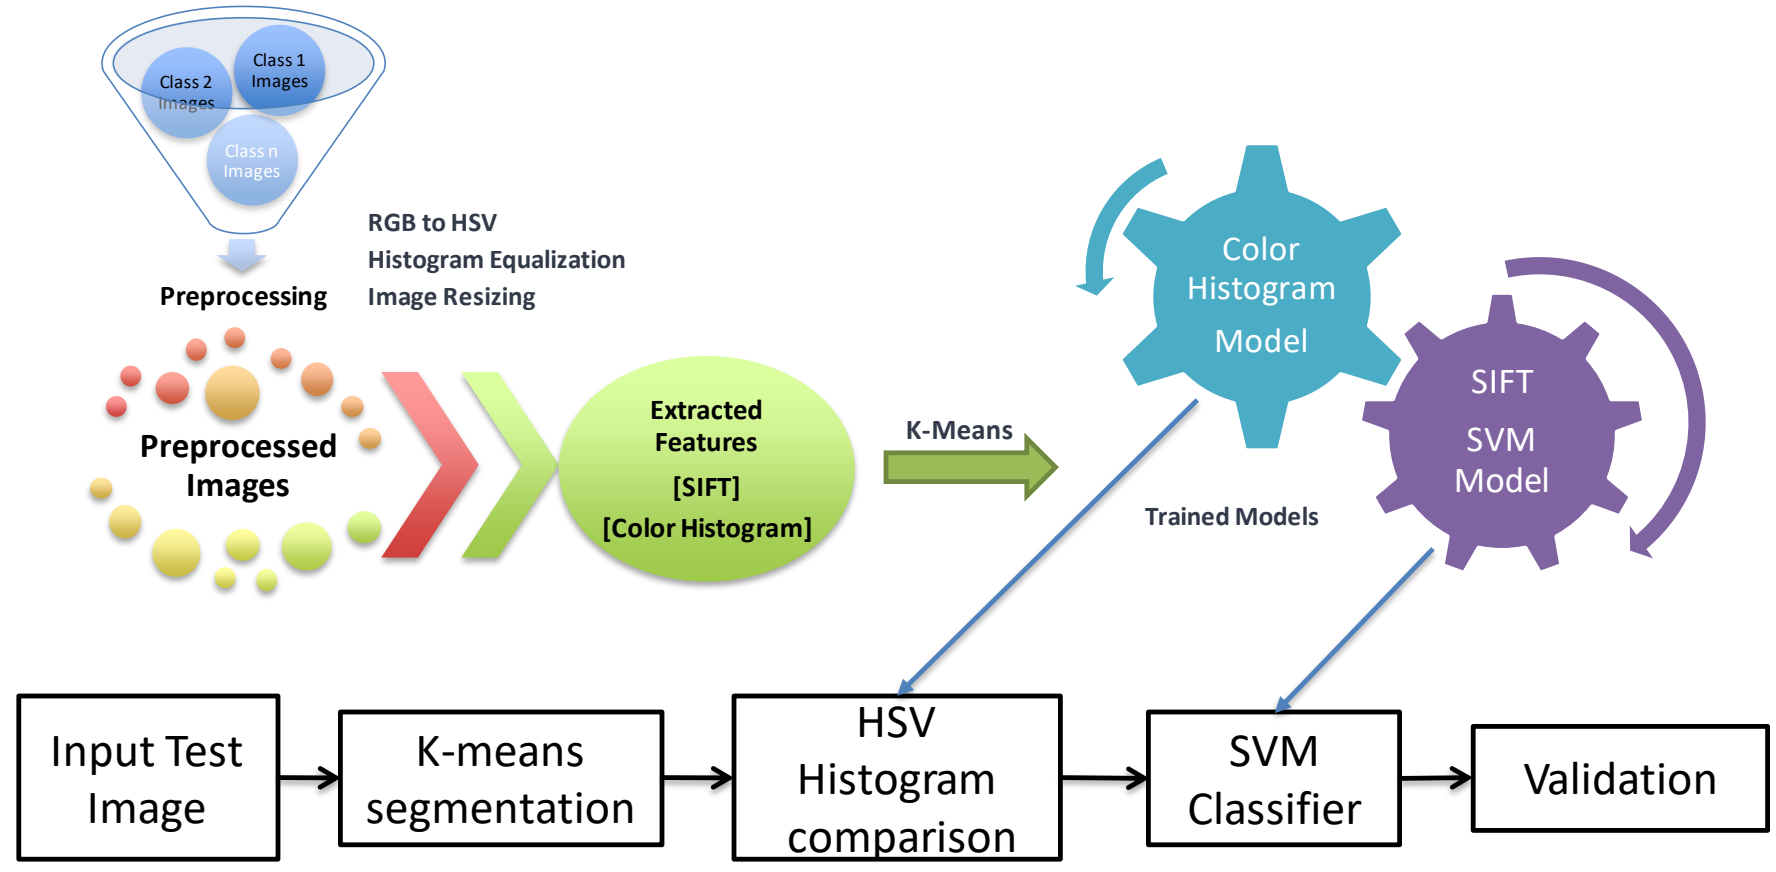

Pre-Processing

- Contrast Enhancement
  - Used HSV over RGB for color histogram
  - Performed histogram equalization on Intensity values
  - RGB to Grayscale for SIFT
- Improving Efficiency
  - Manual segmentation of training images to focus on regions of interest
  - Images resized such that max(row, col) = 1000 pixels

Color Histogram

Segment test image

↓

Get H,S,V for foreground segment

↓

Compare H,S, V histograms of test image to trained histograms

- Use Manhattan distance as comparison
- Combine Manhattan distance and weigh H more
- Rank distances by class

↓

Outputs rank of class histogram similarity

Training – Color Histogram

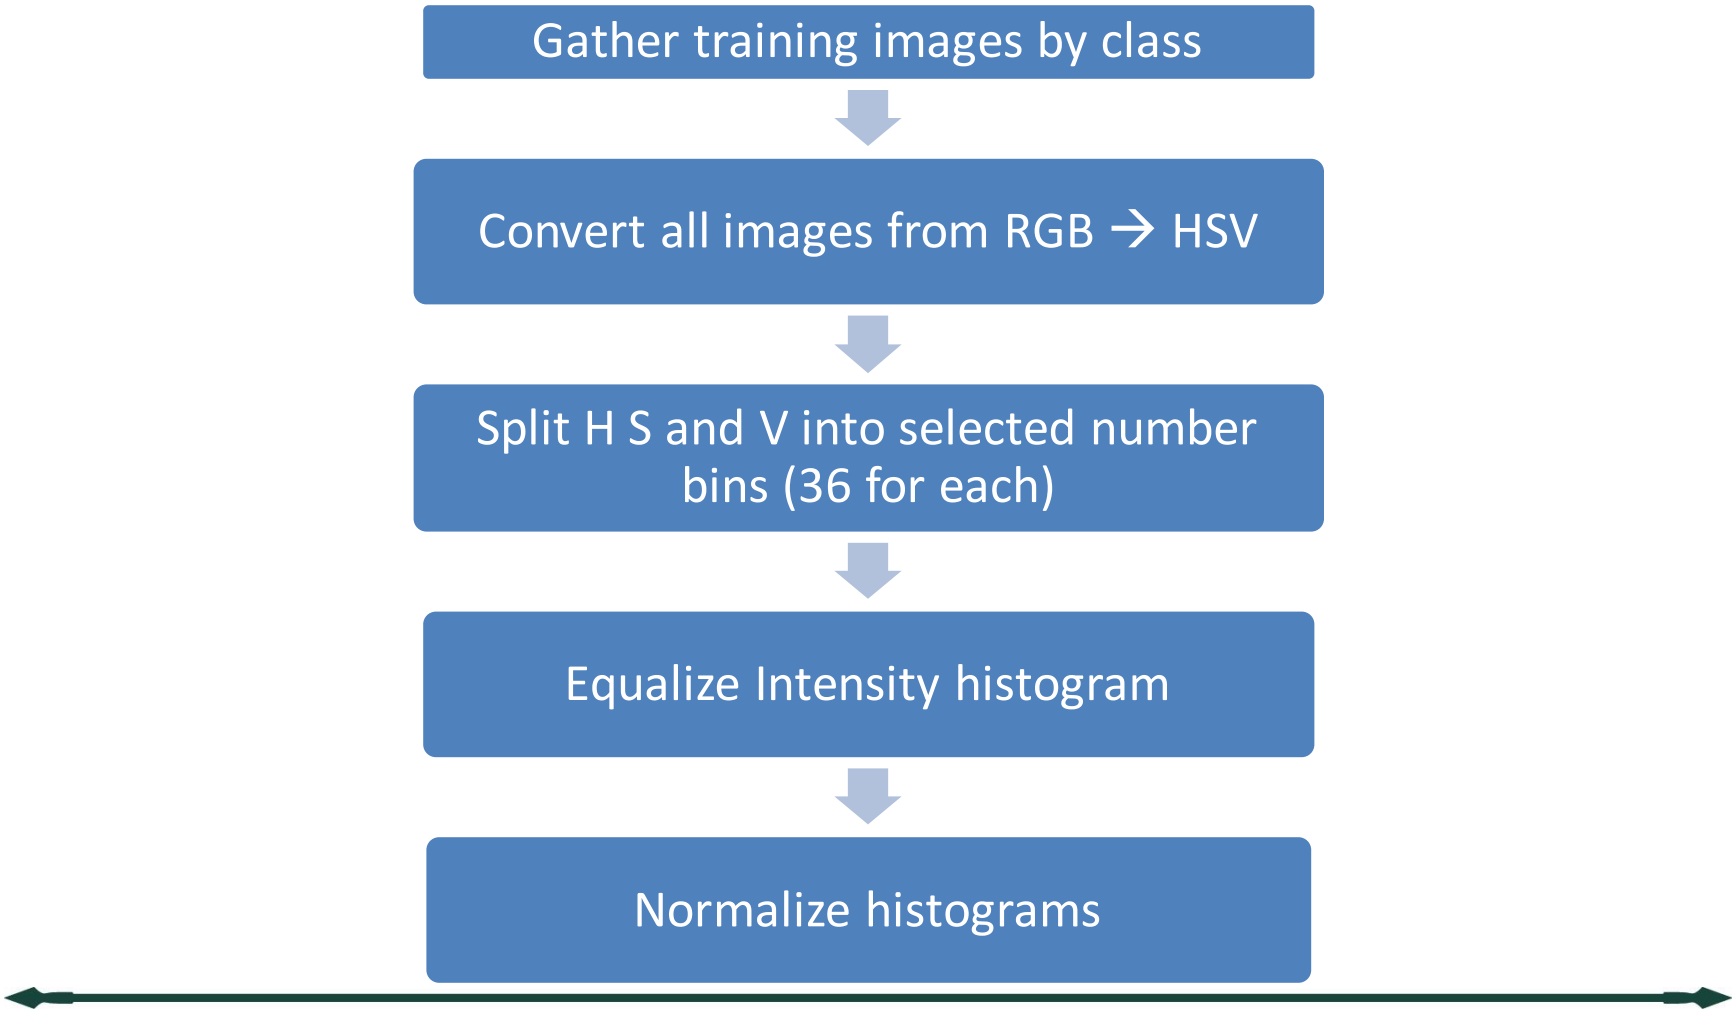

Feature Extraction

SIFT : Scale Invariant Feature Transform

For each image,

- Extract SIFT: detects local key points (N) at different scale
- Returns N x 128 key point descriptors with their locations
  - How to make it a global descriptor?

Apply K-Means..

- Tried different values of K, worked best for K=10
- Total Clusters for each class = 50 x 10 = 500 key point clusters
- 10 Sift Models for each class with 500 x 128 feature vectors

Training – SIFT

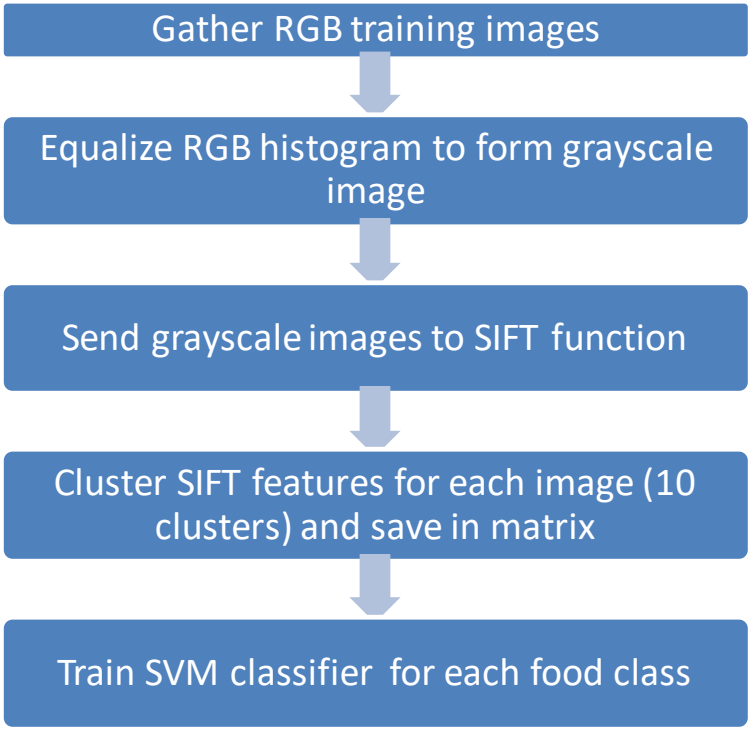

Feature matrix passed to SVM is a 600 row matrix

- First 300 rows are SIFT features belonging to food class being trained
- Last 300 rows are SIFT features belonging to other classes
- Features randomly chosen

Optimization By Filtering SIFT Features: Utensil Class

- Analyzed SIFT key points
- Use knowledge about most common backgrounds? Utensils?
- Extracted sift features for a new negative class “Utensils”
- 40 images consisting of bowls, plates, baskets in different view-angles
- Improvement of 3% in Rank-1 accuracy






Results and Critique

Results

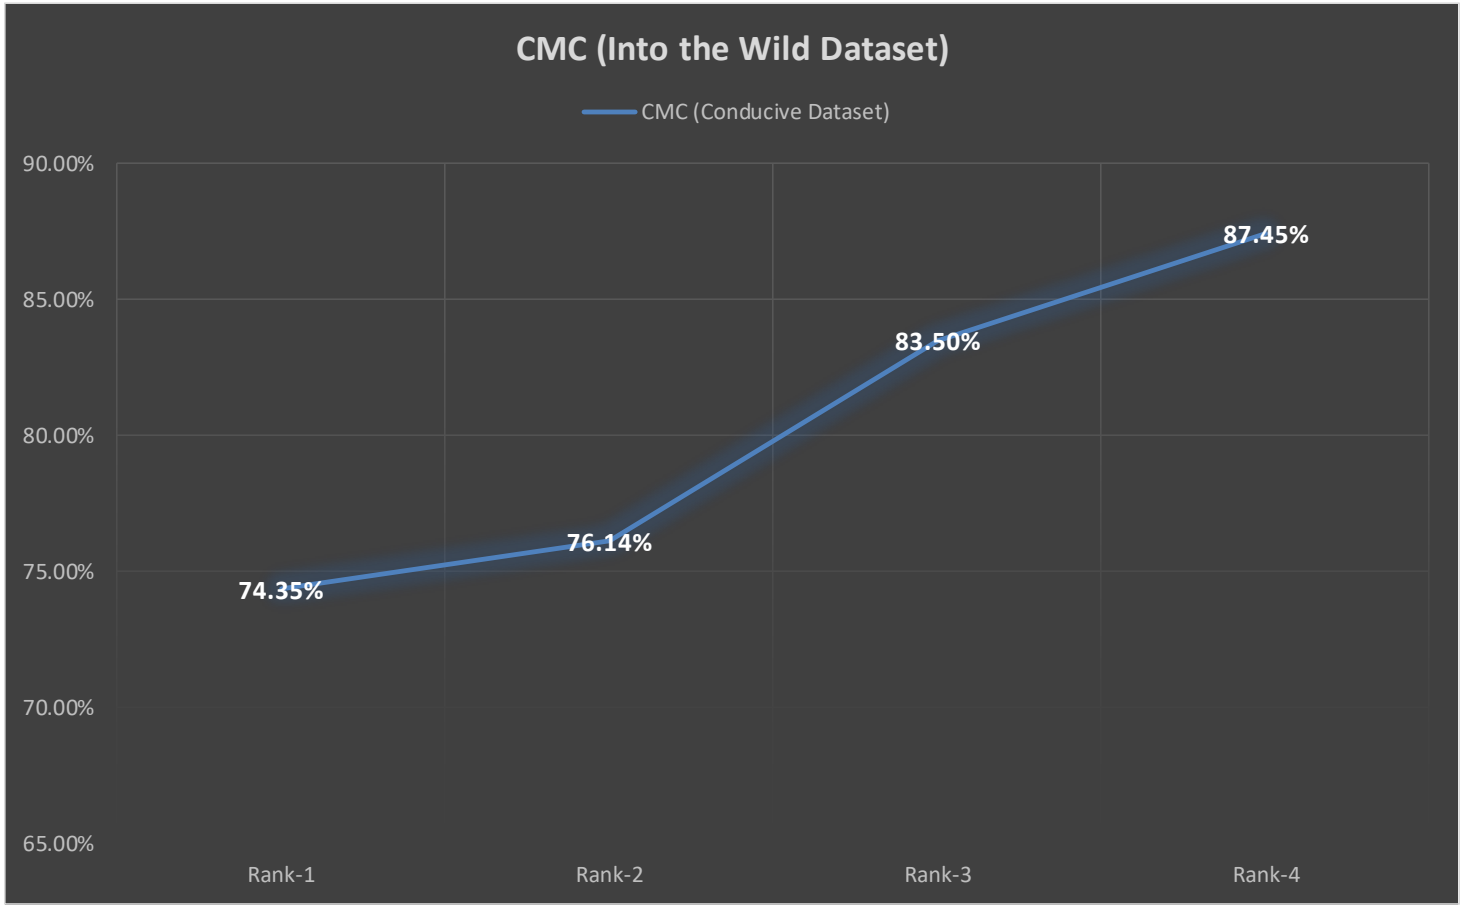

Time Efficiency : 89.40 seconds

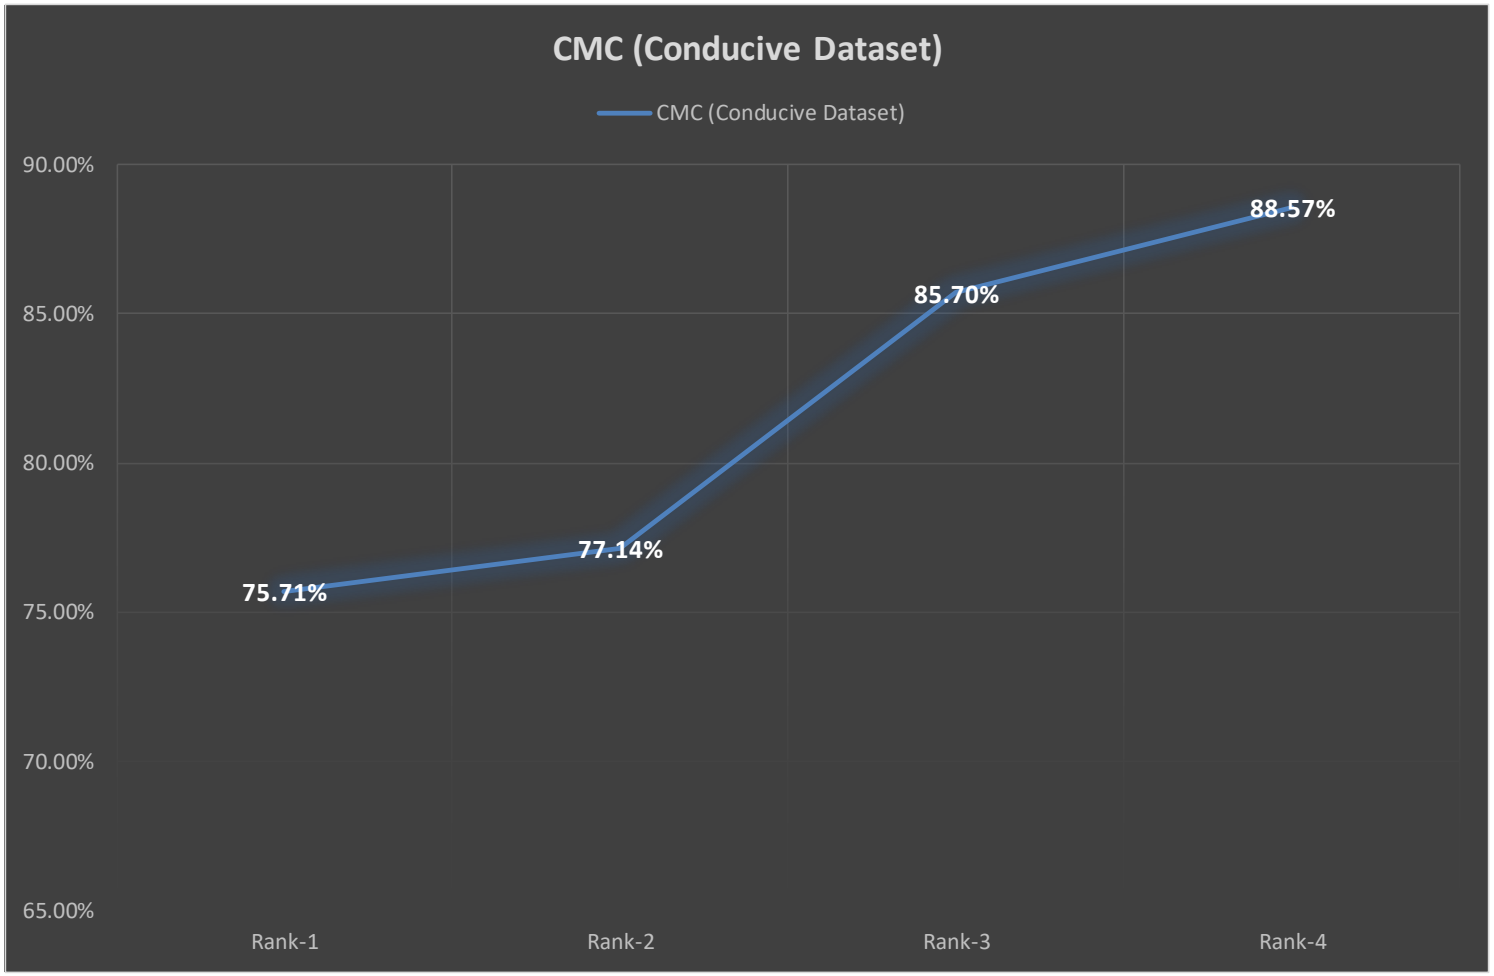

Time Efficiency : 88.40 seconds

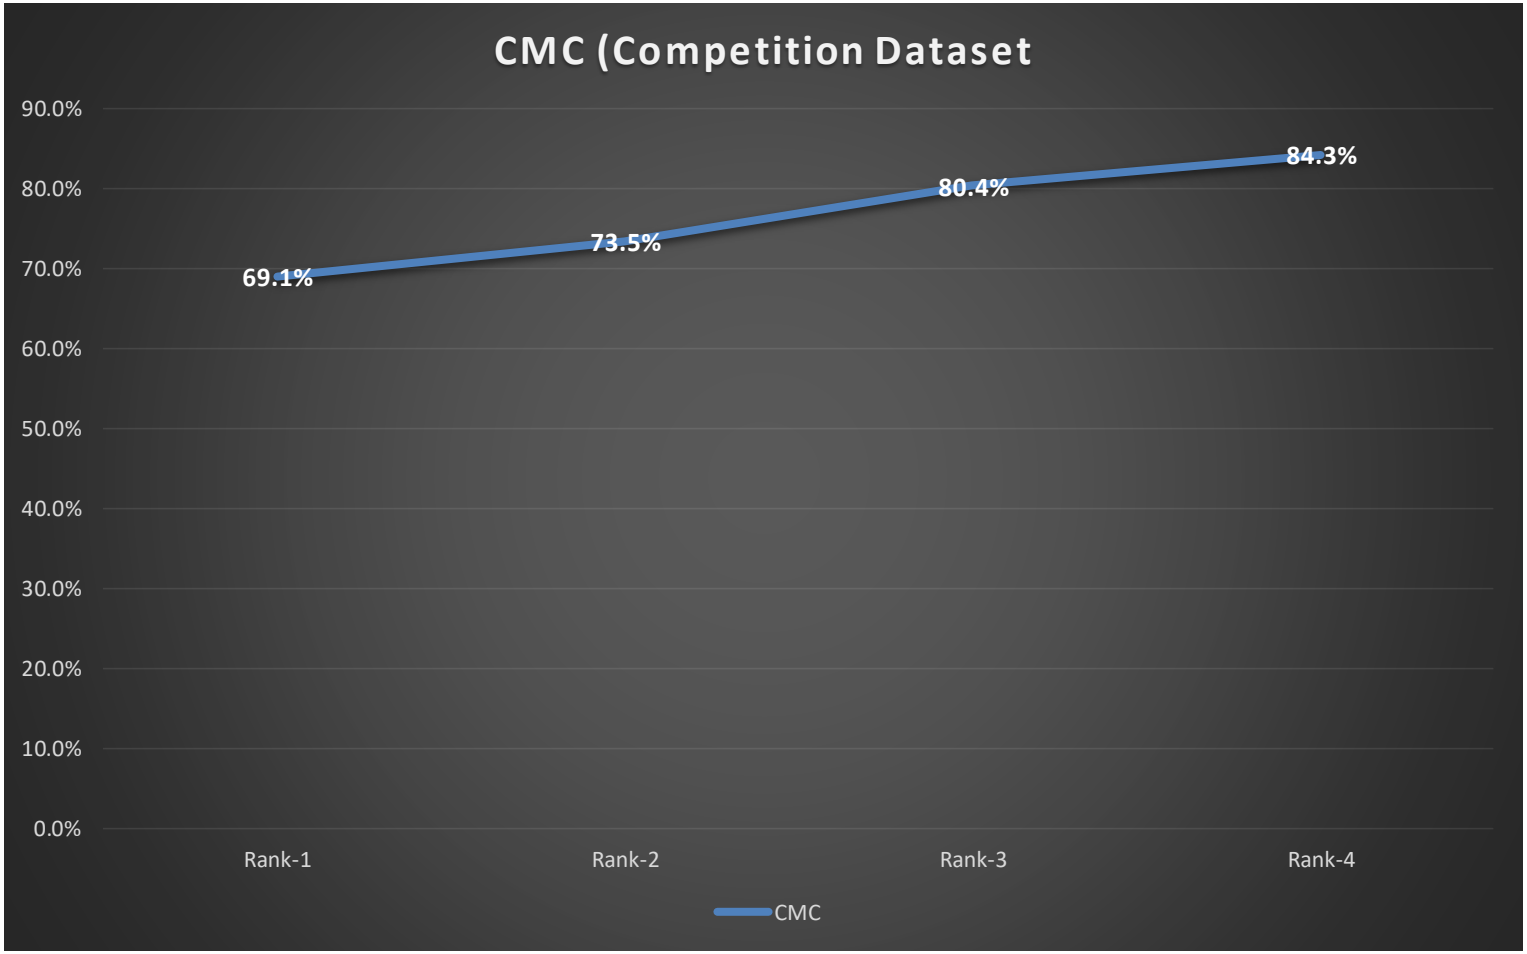

Time Efficiency : 57.72 seconds

Discussion of Algorithm

- Color was by far best classifier
- Features attempted but not used:
  - Watershed segmentation
  - Histogram of Oriented Gradient
- Some classes are bound to be falsely classified
  - Apple, Tomato (color)
  - French Fries, Banana
  - Hotdog, Banana (shape)

Lessons Learnt

- Very hard to balance between accuracy and time requirement / computational efficiency of the system
- A very thin line between generalization and over-fitting of the model
- Segmentation is a big challenge
- More the training dataset, better it is
- Need to think intuitively about the relations/ patterns that exist and suitably optimize the parameters of various computer vision modules (feature extraction, training classifier, etc.)

Future Work

- To explore the Deep Neural Networks to identify the hidden relations between the images of each class
- Use of gradient features (eg. HOG) in conjunction to color histograms
- Building an android app for this task
- Estimating the quantity of the food by predicting the size of any known objects in the background
- Use of 3D alignment and matching to predict the distance of food from camera

References

- [1] http://www.cse.msu.edu/~cse803
- [2] S. Yang, M. Chen, D. Pomerleau, and R. Sukthankar, “Food recognition using statistics of pairwise local features,” in Proc. of IEEE Computer Vision and Pattern Recognition, 2010.
- [3] Z. Zong, D.T. Nguyen, P. Ogunbona, and W. Li, “On the combination of local texture and global structure for food classification,” in Proc. of International Symposium on Multimedia, pp. 204–211, 2010.
- [4] M. Chen, K. Dhingra, W. Wu, L. Yang, R. Sukthankar, and J. Yang, “PFID: Pittsburgh fast-food image dataset,” in Proc. of IEEE International Conference on Image Processing, pp. 289–292, 2009.
- [5] www.google.com
- [6] www.bing.com
- [7] D. G. Lowe. Distinctive image features from scale-invariant keypoints. IJCV, 60(2):91–110, 2004.
- [8] http://www.mathworks.com/help/stats/kmeans.html
- [9] CORTES, C. AND VAPNIK, V. 1995. Support-vector network. Mach. Learn. 20, 273–297.

Questions?